

# Performance Analysis Model

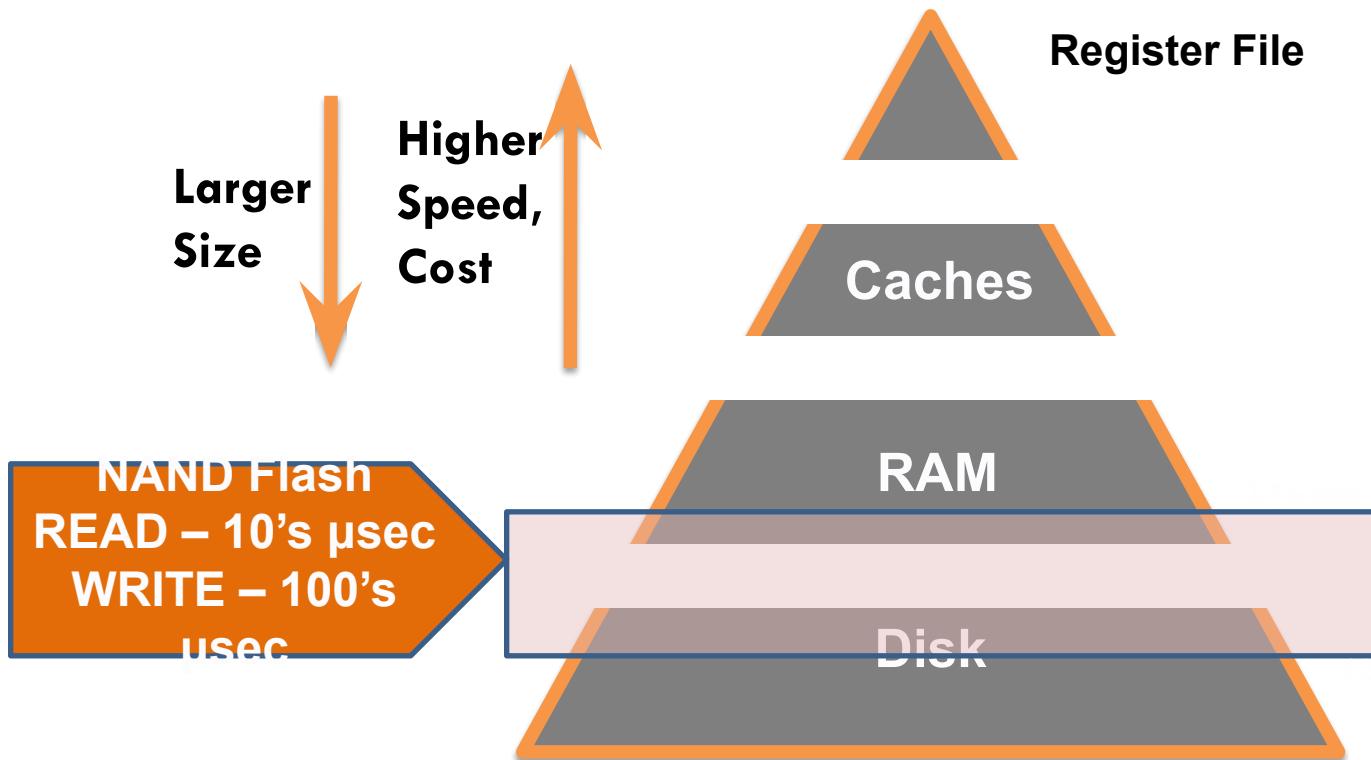
Jihong Kim

Dept. of CSE, SNU

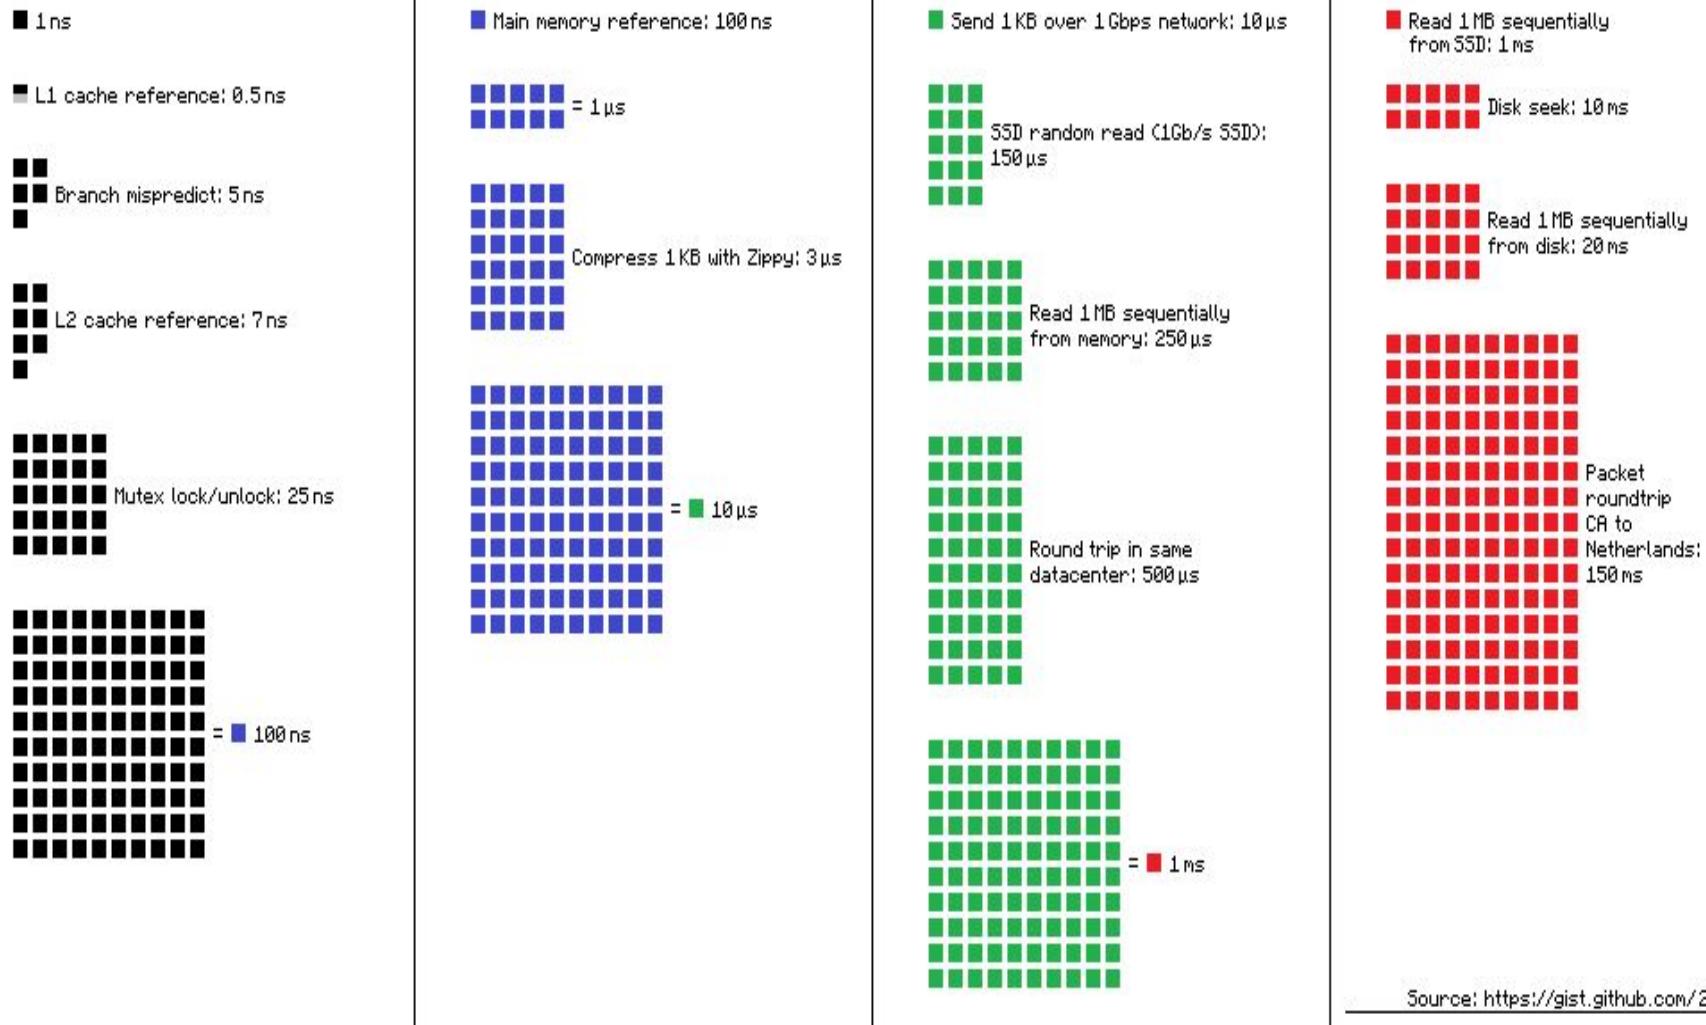
# Performance Metrics

- Throughput
- Response Time
- Capacity
- Reliability
- Cost
- ...
- Used for CPU, Memory, SSD, Network, ....

# Memory Hierarchy



# Latency Numbers Every Programmer Should Know



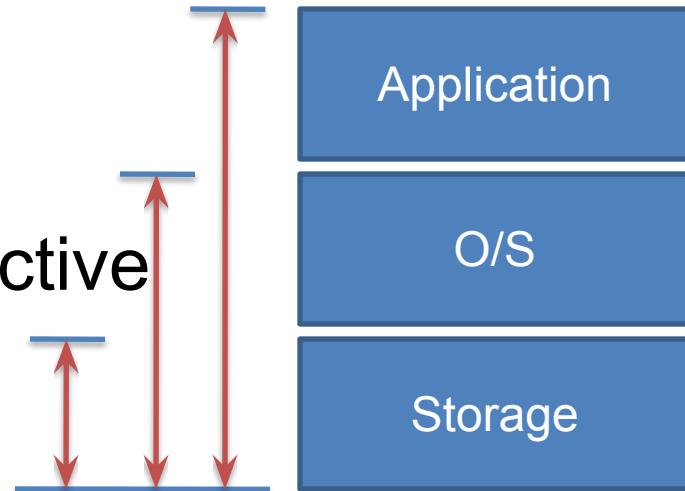
Source: <https://gist.github.com/2841832>

# Throughput

- IO rate
  - IOPS (accesses/second)
  - Used for applications where the size of each request is small
  - e.g. transaction processing
- Data rate
  - MB/Sec or bytes/Sec
  - Used for applications where the size of each request is large
  - e.g. scientific applications

# Response Time

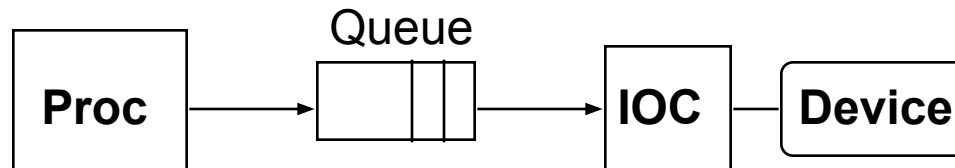
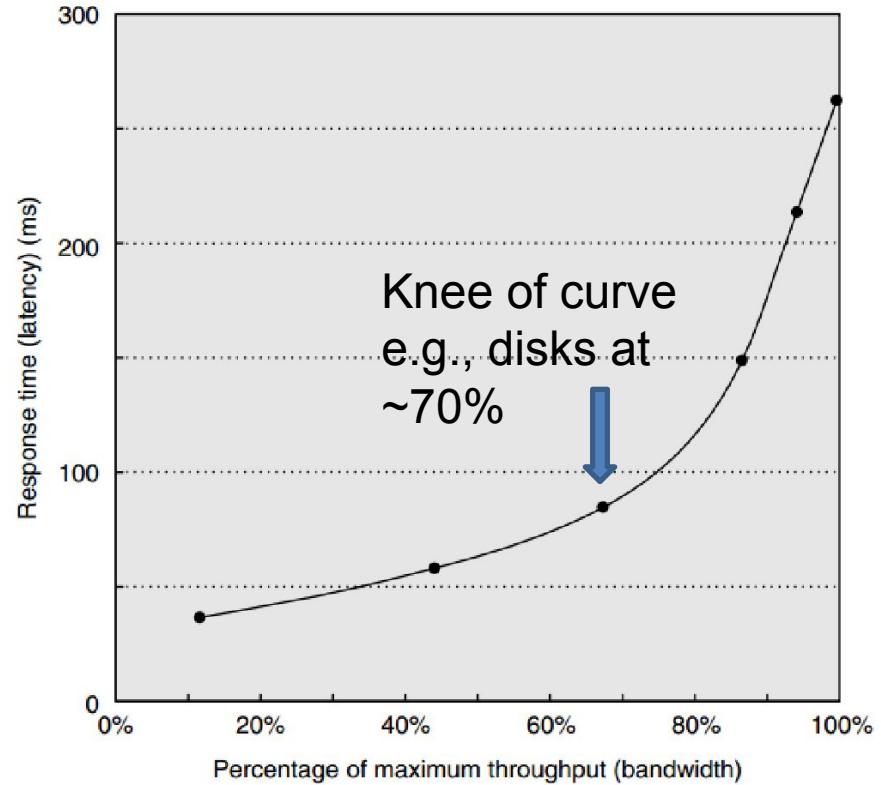
- How long a storage system takes to access data
- Depending on what you view as the storage system
  - User's perspective
  - OS's perspective
  - Disk controller's perspective



# Disk I/O Performance

**Metrics:**  
**Response Time**  
**Throughput**

**Q: how to balance  
response time &  
throughput?  
(e.g., adding more servers)**



**Response time = Queue + Device Service time**

# Performance Analytic Model

- Better to understand key performance factors
- High-level insight on system performance
- Complementary to other performance evaluation methods:
  - Measurement-based
  - Simulation-based

# Example Analytic Model

- Original Amdahl's Law
  - Useful in understanding an upper bound on overall speedup (insight with less accuracy!!)

S: speedup

$$S = \frac{t_p + t_s}{\frac{t_p}{N} + t_s},$$

N: no. of processors

f: fraction of sequential part

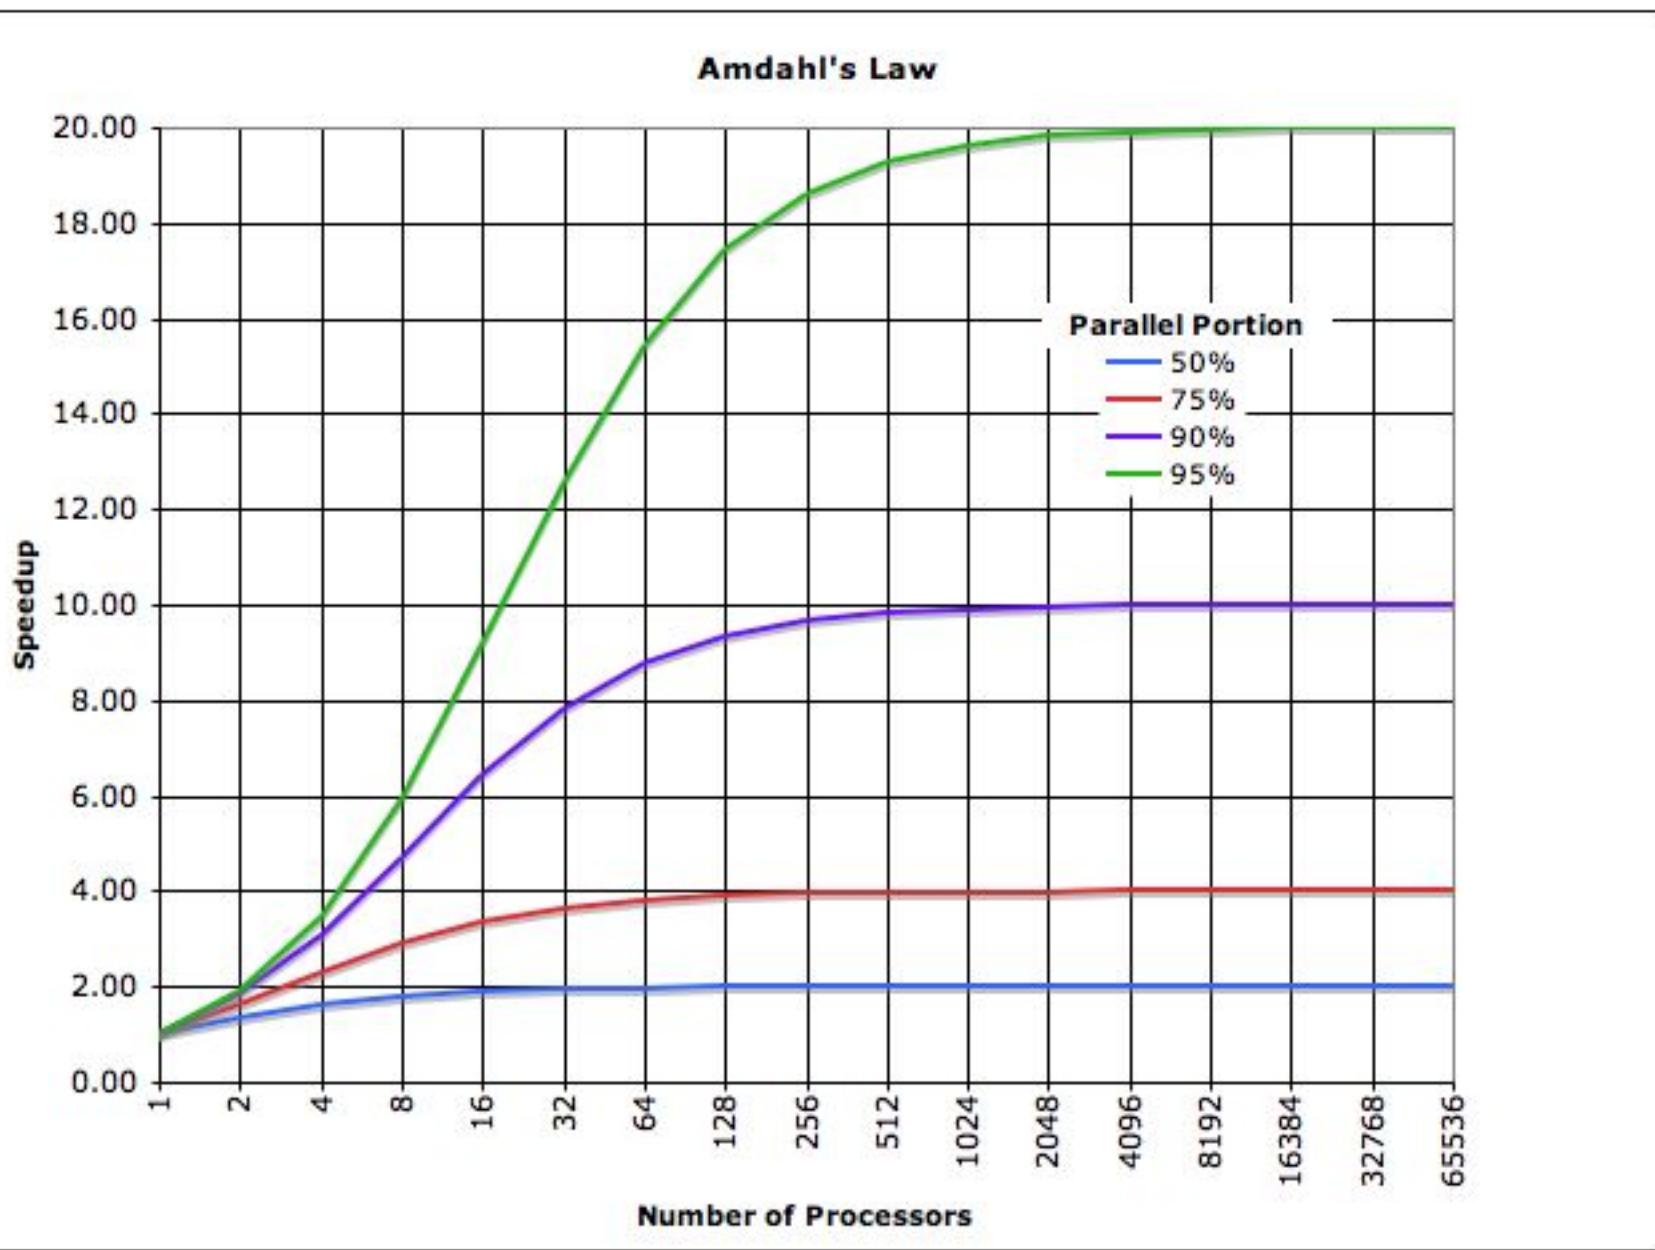
$t_p$  and  $t_s$ :

exec. time for parallel  
and sequential part

# **Example: 100 processors**

$$f = \frac{S}{N} \cdot \frac{1}{1 - \frac{1}{N}}$$

- **S = 100**
  - **f = ?**
- **S = 99**
  - **f = 1/(99\*99) = 0.000102**



# Amdahl's law (in general)

**Execution Time** *after improvement*

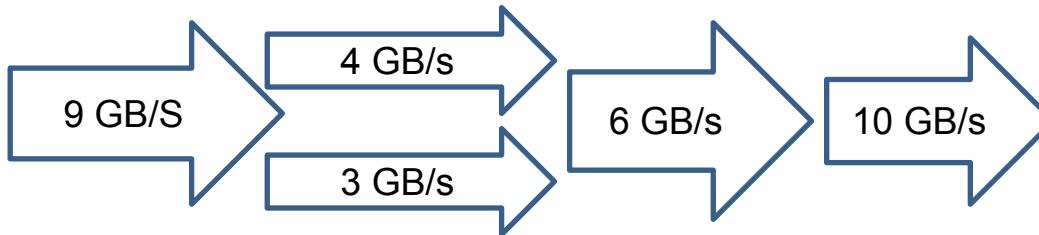
$$= \frac{\text{Execution Time}_{aff\epsilon}}{\text{Amount of i}}$$

+ **Execution Time** *unaffected by improvement*

# Bottleneck Analysis

- An upper bound on the throughput of a system from the throughput of its subsystems.
  1. The max. throughput of K subsystems **in parallel**  
= the **sum** of the subsystem throughputs.
  2. The max. throughput of K subsystems **in series**  
= the **minimum** of the subsystem throughputs.

# Example: Bottleneck Analysis



- Max Bandwidth =  $\text{MIN}(9, (4+3), 6, 10) = 6$
- The most expensive subsystem should be bottlenecked.
- Non-bottlenecked throughputs should be decreased to save cost.

# Roofline Model

- A visual way to identify performance bottlenecks and to find potential optimization directions
- Bottleneck analysis between application characteristics and architectural capabilities

$$FLOPS = \frac{\#FLOP}{sec} =$$

$$= \text{Arithmetic Intensity} \times \text{Bandwidth}$$

$$= AI \times BW$$

$$\text{Attainable FLOPS} = \min(\text{AI} \times \text{BW}, \text{Peak FLOPS})$$



$$\text{Arithmetic Intensity} = \frac{\#FLOP}{\text{byte}}$$

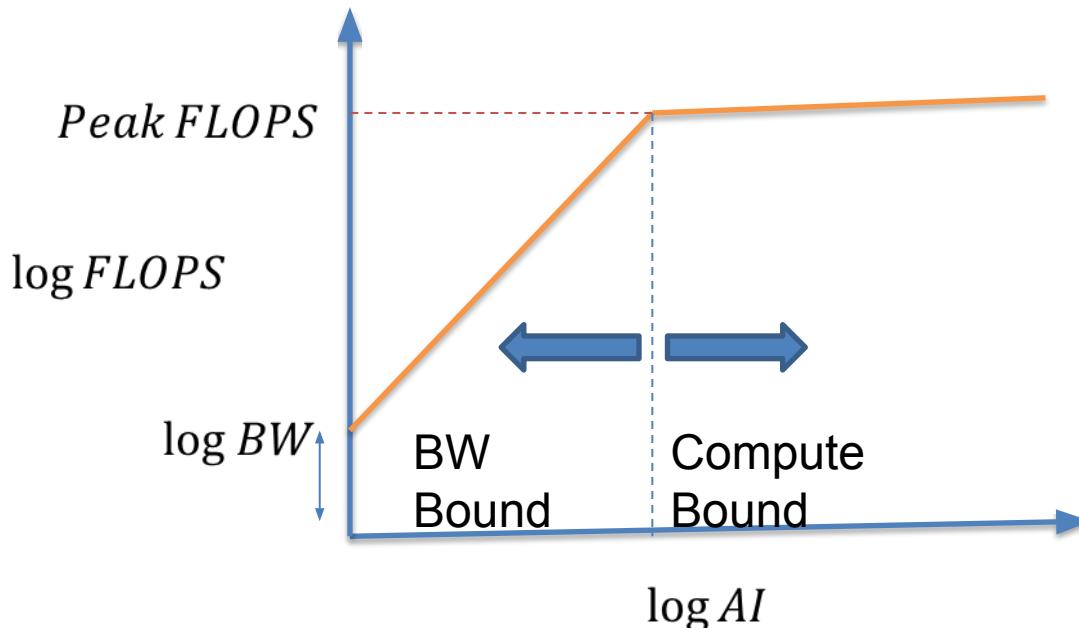
- The number of operations per unit data
- Algorithm specific
- Example AI:
  - Matrix Addition of two  $N \times N$  matrices
    - Total number of data accesses =  $2NN$  (reads) +  $NN$  (writes) =  $3NN$
    - Total number of operations =  $NN$  (adds)
    - $AI = \frac{NN}{3NN} = \frac{1}{3}$  (constant)
  - Matrix Multiplication of two  $N \times N$  matrices
    - $AI = \frac{(2N-1) \times NN}{3NN} = \frac{2N-1}{3}$  (as  $N$  gets higher, AI increases)

# Roofline Graph

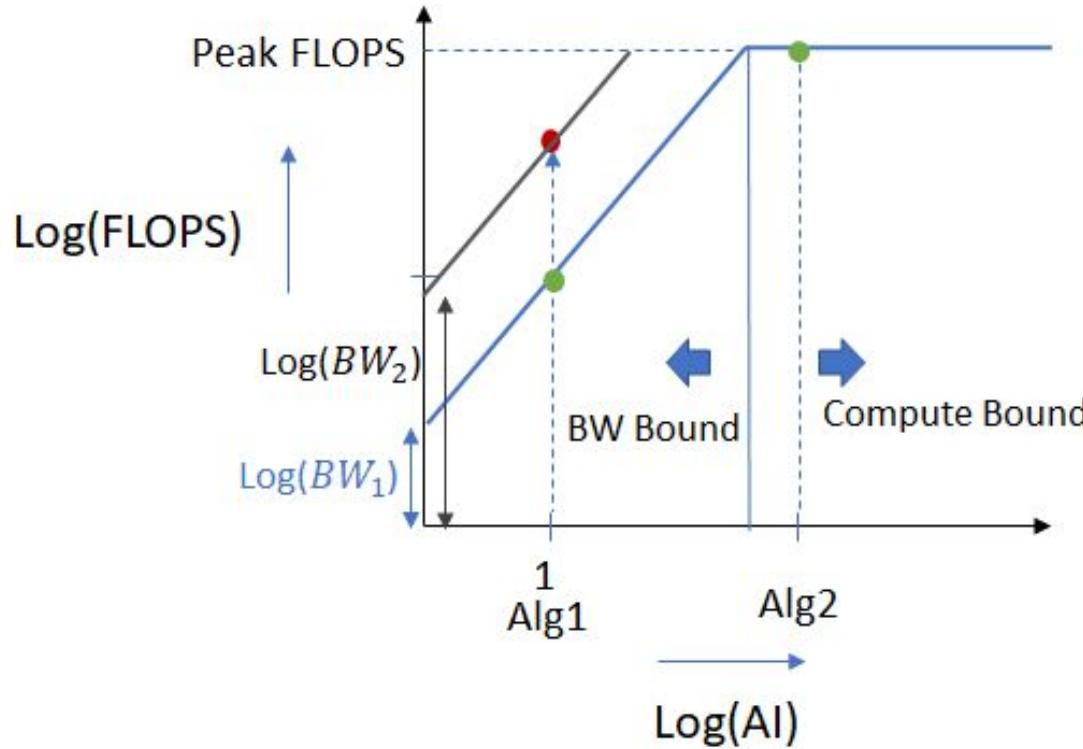
- $FLOPS = AI \times BW$

$$\log FLOPS = \log AI + \log BW$$

$$Y = X + C$$

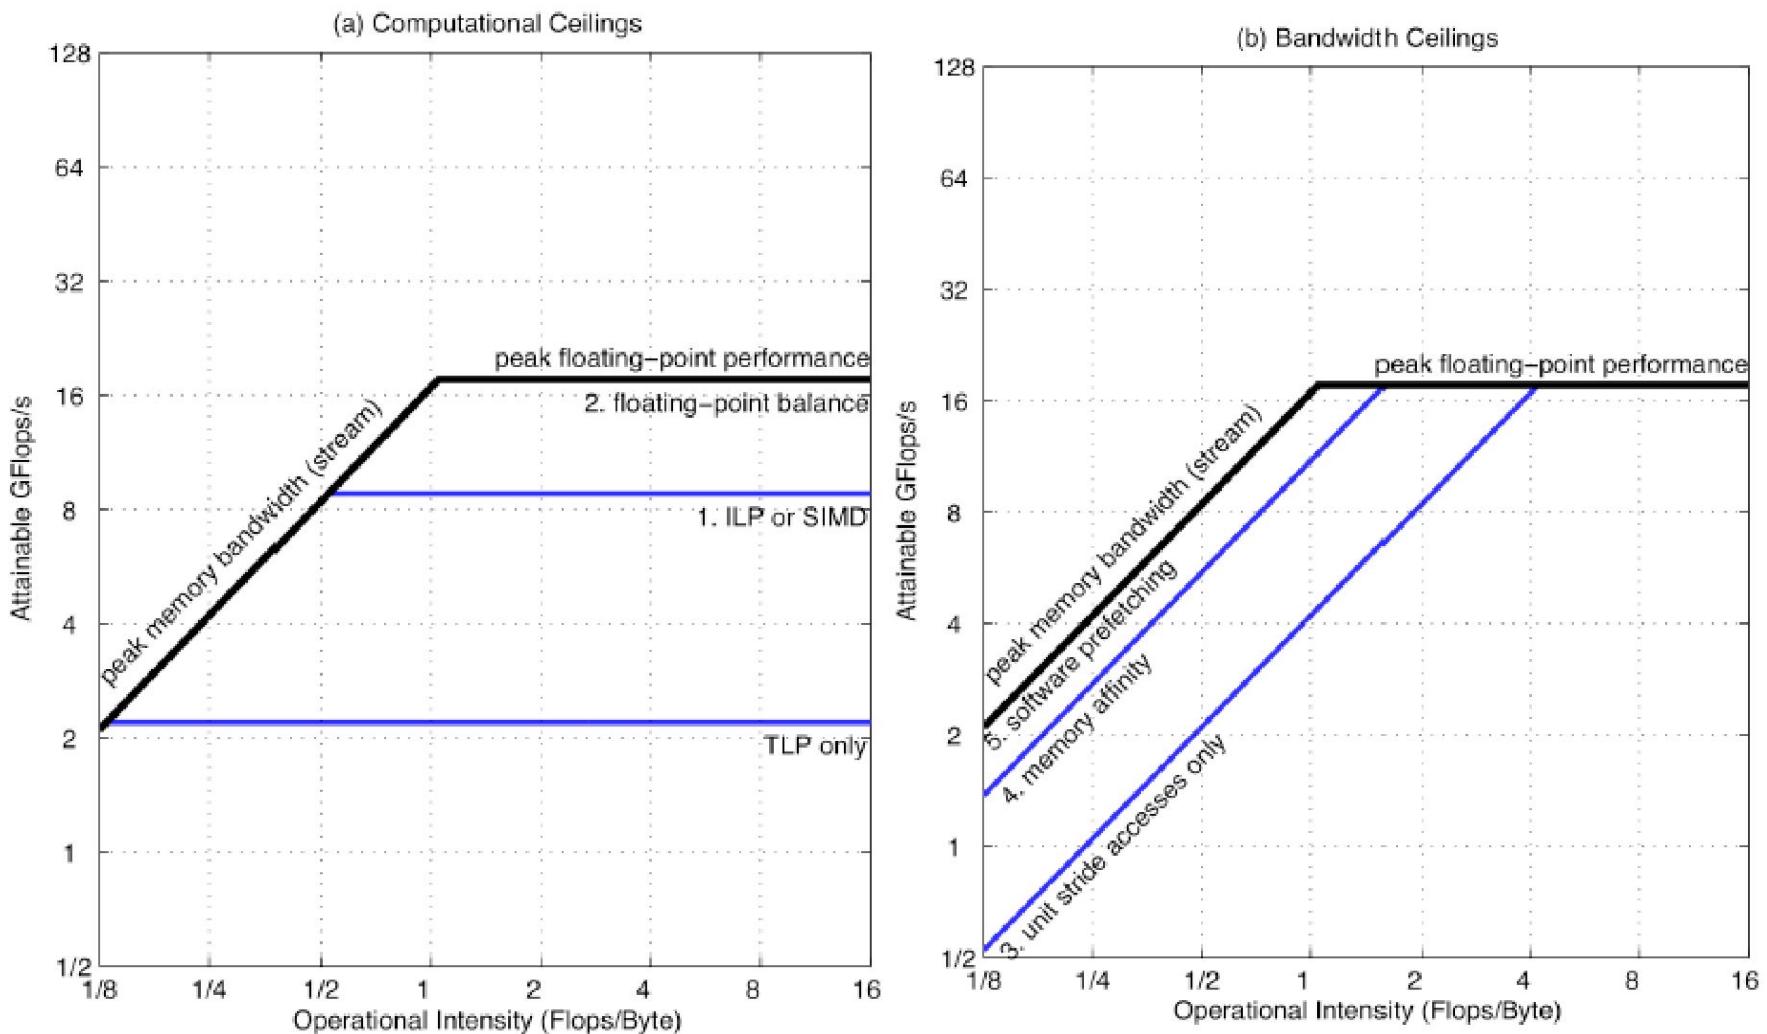


# Optimization Guidelines



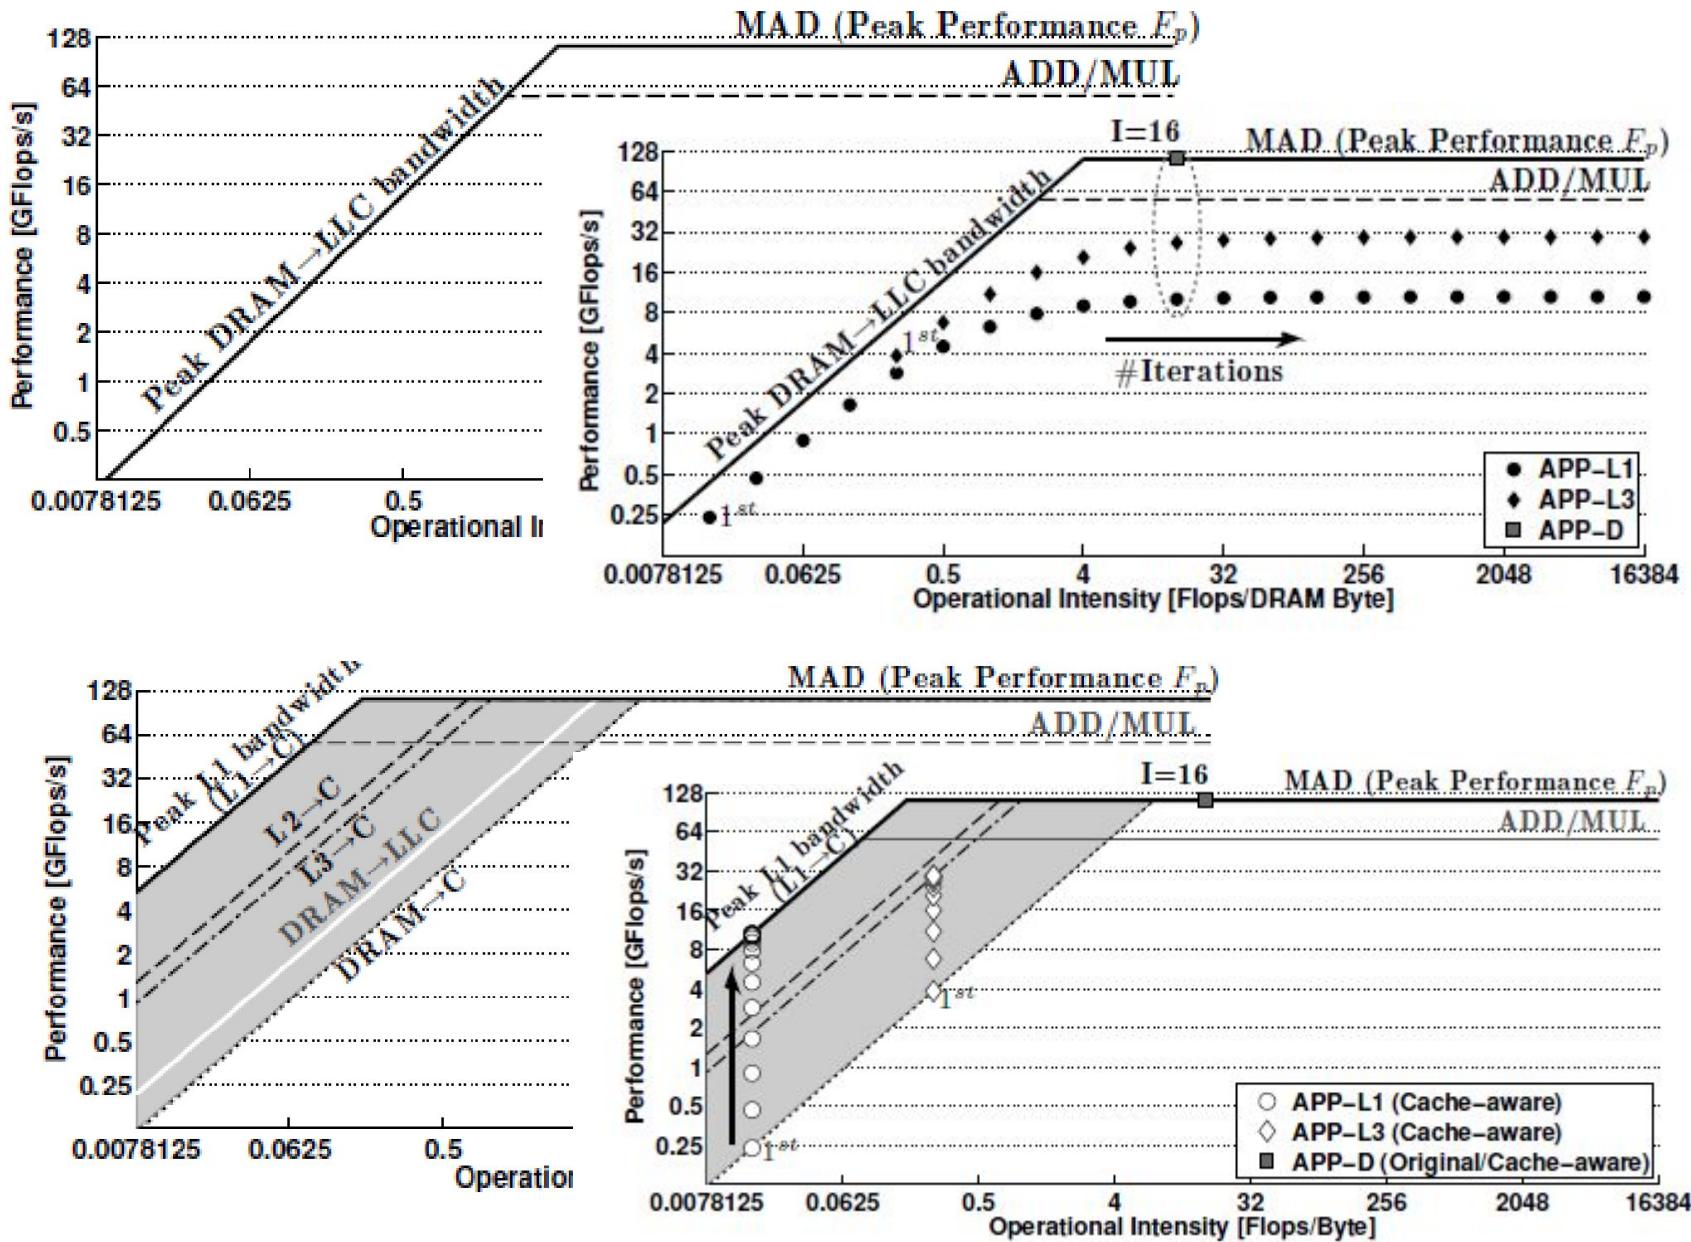
- Alg 1: Move Right (higher data reuse)  
Move Up (higher BW memory)
- Alg 2: Move Up (more powerful CPU)

# Roofline Model w/ Ceilings



# Cache-Aware Roofline Model

- Original Roofline Model
  - AI varies for a given algorithm:  
(e.g., different loop iterations or different optimizations)
- Cache-aware Roofline Model
  - AI is constant.  
□ AI is constant.



# APP-L1 vs. App-L3 vs. APP-D

- 정의: 같은 데이터셋을 N회 반복. 첫 반복만 DRAM 접근, 이후는 각 클래스의 주 접근 레벨로 동작.
  - APP-L1: 첫 반복 이후 주 접근 = L1
  - APP-L3: 첫 반복 이후 주 접근 = L3
  - APP-D: 모든 반복에서 DRAM 접근

# Table: 1회 반복당 지표 ( $\beta_0$ , $\varphi_0$ , I)

클래스	주 접근 레벨 (반복 이후)	$\beta_0$ [bytes]	$\varphi_0$ [flops]	I (첫 반복)
APP-L1	L1	4.5 KB	72	$1/64 \approx 0.0156$
APP-L3	L3	2 MB	524.3 K	$1/4 = 0.25$
APP-D	DRAM	61 MB	1023.4 M	16

# Cache-Aware에서는 왜 I(작업강도)가 고정되나?

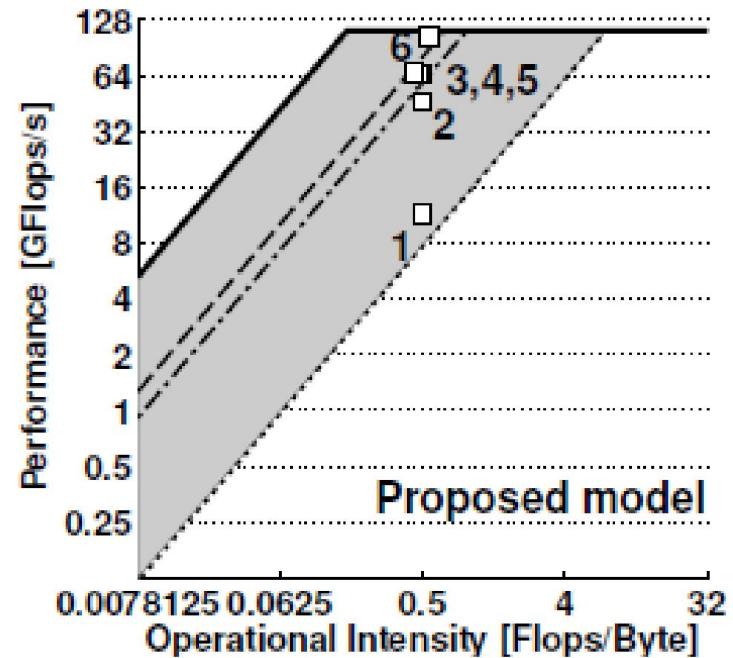
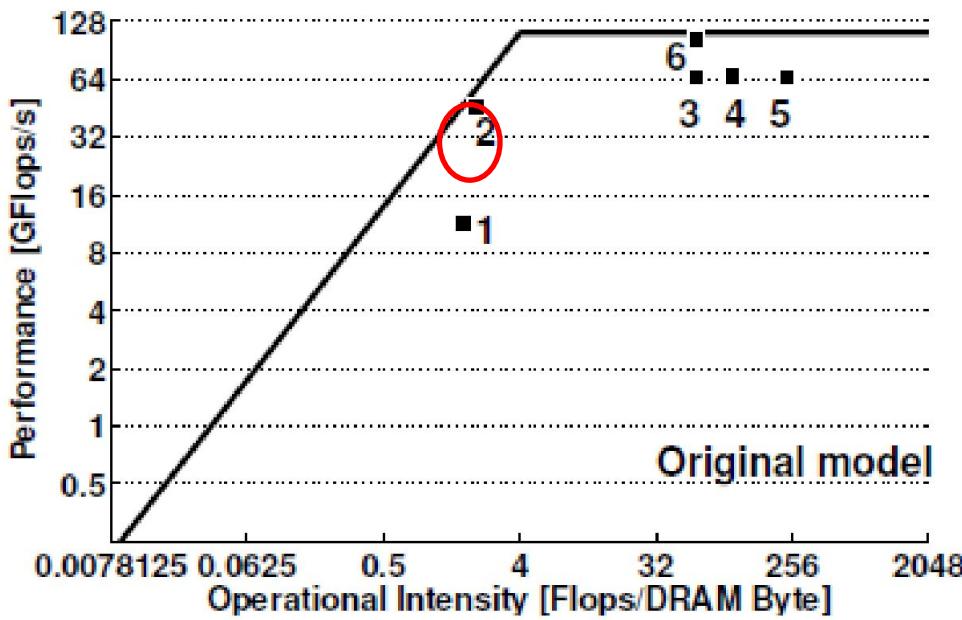
- 정의:  $I = \varphi / \beta = (\varphi_0 \times N) / (\beta_0 \times N) = \varphi_0 / \beta_0 \rightarrow N$ 이 소거되어 상수
- APP-L1, APP-L3, APP-D 모두 Cache-aware 정의에선 반복이 늘어도 같은 I를 유지
- Original DRAM- 중심 모델에서는 L1/L3의 경우 DRAM 바이트가 고정되어 I가 오른쪽으로 이동

# 반복 N 증가에 따른 Flops/S 변화

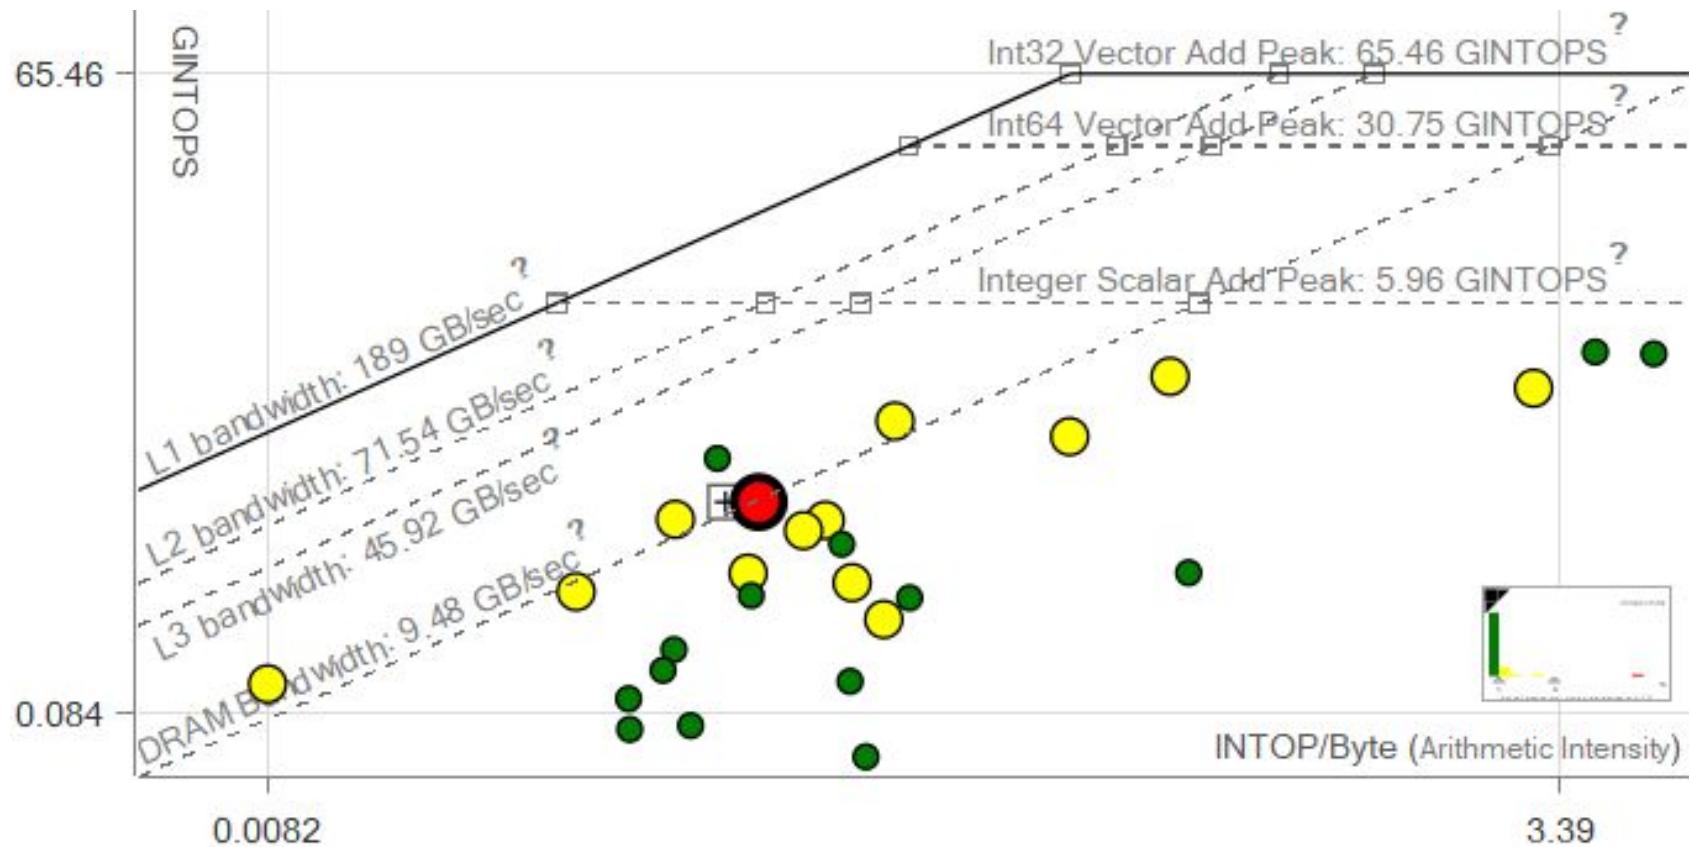
- APP-L1: 첫 반복 DRAM 접근의 지연이 평균에서 희석 → L1→Core 대역폭 천장으로 수렴하여 Flops/s↑
- APP-L3: 첫 반복 후 L3 히트가 지배 → L3→Core 천장으로 수렴하여 Flops/s↑
- APP-D: 매 반복 DRAM 접근 → DRAM→Core 경로에 묶여 상대적으로 평坦

# Cache-aware vs Original Roofline

- In the original roofline model, if an application is in the memory bound region, it is difficult to achieve the maximum attainable performance.

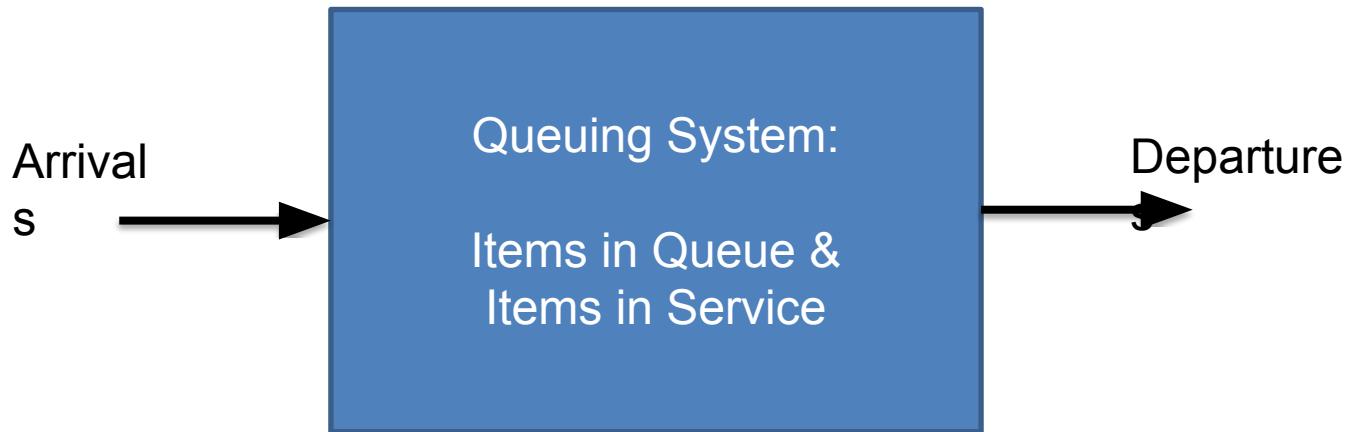


# Cache-Aware Roofline Model in Intel Advisor



- Large red dot: takes up the most time,
  - candidate (loop or function) for optimization

# Little's Law: $L = \lambda W$



$L$  = average number of items in the queuing system

$W$  = average waiting time in the system for an item (i.e. latency)

$\lambda$  = average number of items arriving per unit time  
(i.e., bandwidth)

**Assumptions** Stable, work-conserving system. Use long-run averages

Good Tool for Meeting Service-Level Objective (SLO) of a system

# Example 1A: 평균 재학기간

- 매년 신입생이 2000명인 학교에서
  - $\lambda = 2000/\text{년}$  (avg throughput)
- 평균적으로 학교를 다니는 학생이 3000명이면
  - $L = 3000$

Q: 학생 1인당 학교에 다니는 기간은?

$$(\text{avg latency}) \quad W = 3000/2000 = 1.5\text{년}$$

## Design knobs via $L = \lambda W$

**Cap L:** 동시 등록한 학생 수를 제한

**Reduce W:** 개설 강좌 증대, 선수과목 축소, 졸업학점

축소

**Bound  $\lambda$ :** 신입생 수 제한

# Example 1B: 평균응답시간

- 매일 평균 50개의 이메일을 받는 사람이
  - $\lambda = 50/\text{day}$  (avg throughput)
- 평균적으로 InBox에 150개의 응답하지 않은 이메일을 가지고 있다면
  - $L = 150$

Q: 평균응답시간은?

- $W = 150/50 = 3\text{일}$

## Design knobs via $L = \lambda W$

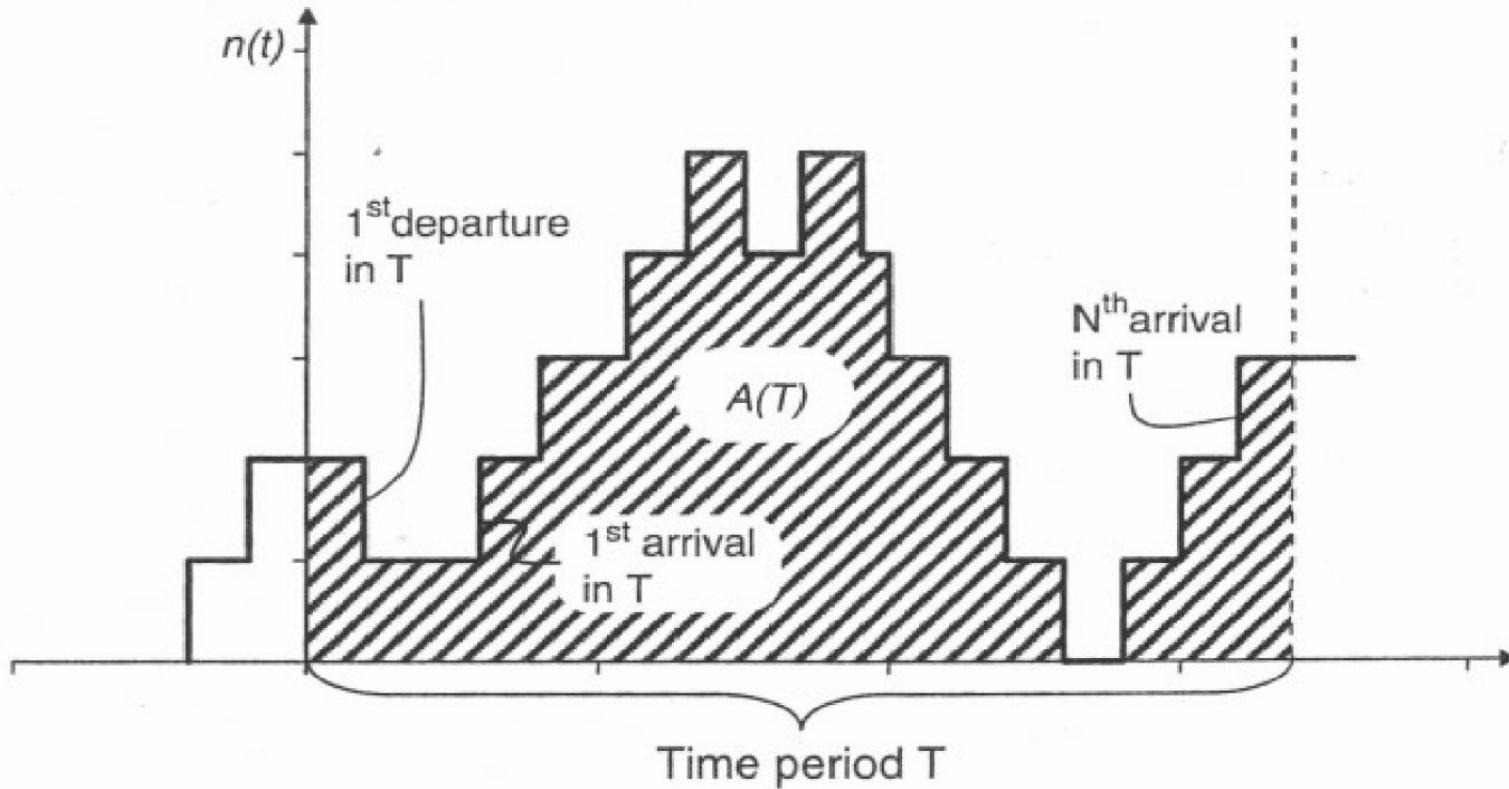
**Cap L:** Set the max. inbox size

**Reduce W:** faster replies using templates

**Bound  $\lambda$ :** rate-limiting policy (office hours

only)

# Intuitive Argument



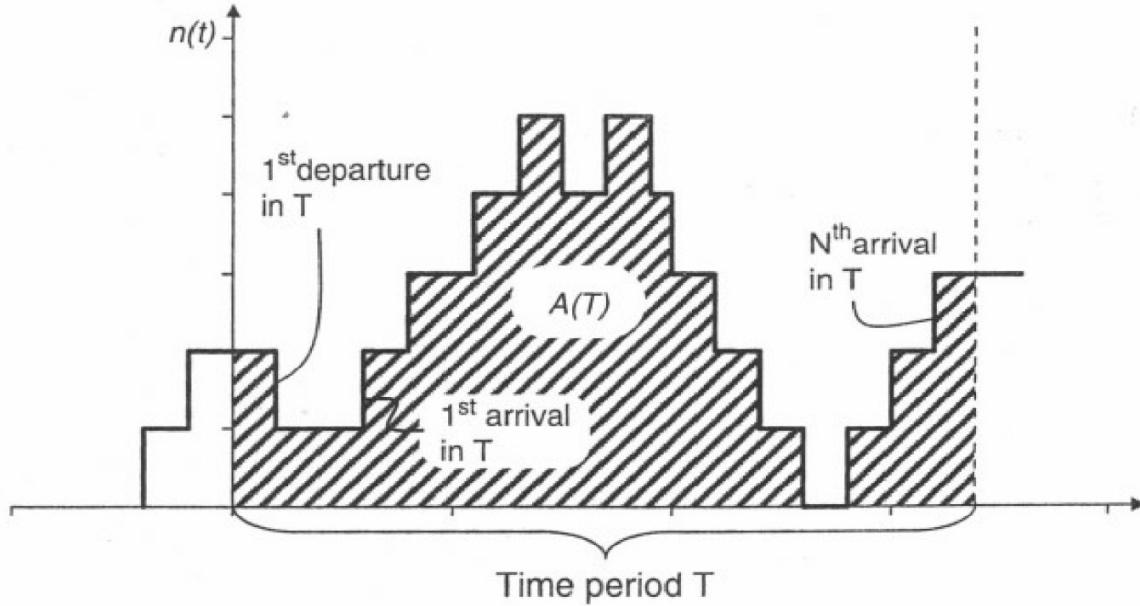
$n(t)$  = the number of items in queue at time  $t$

$T$  = a long period of time

$A(T)$  = the area under  $n(t)$  over  $T$

$N(T)$  = the number of arrivals in  $T$

- $\lambda(T) = N(T)/T$
- $L(T) = A(T)/T$
- $W(T) = A(T)/N(T)$
- $\rightarrow L(T) = \lambda(T)W(T)$



$$\lim_{T \rightarrow \infty} L(T) = L$$

$$\lim_{T \rightarrow \infty} \lambda(T) = \lambda$$

$$\lim_{T \rightarrow \infty} W(T) = W$$

# Example 2: Disk Utilization

For a single disk,

- every second, 50 I/O requests are coming in &  
avg disk service time = 10 ms,  
what is the disk utilization?

$$\begin{aligned}\text{U} &: 50 \text{ reqs/s} \times 0.01 \text{ s/req} \\ &= 0.50 \text{ (utilization)}\end{aligned}$$

- For a **stable, work-conserving single server** with **no losses**, the **utilization** (busy fraction) satisfies

- $\rho = \lambda S$ ,

(cf) Little's Law:

- If  $W \leq 0.1 \text{ s}$ ,  $L_{\max} = \lambda \times W_{\max} = 50 \text{ reqs/s} \times 0.1 \text{ s} = 5 \text{ reqs}$

# Example 3: Buffer Requirements

Consider a cache memory with

the cache miss ratio = 6.25% &  
the avg miss latency = 100 ns.

Q: If this cache receives 2 memory references per cycle at 2.5 GHz, how many buffers are needed for recording outstanding misses?

답: avg throughput =

$$\begin{aligned} \text{2 refs/cycle} \times 2.5\text{G cycles/s} \times 0.0625 \text{ misses/ref} \\ = 0.3125 \text{ Gmisses/s} \end{aligned}$$

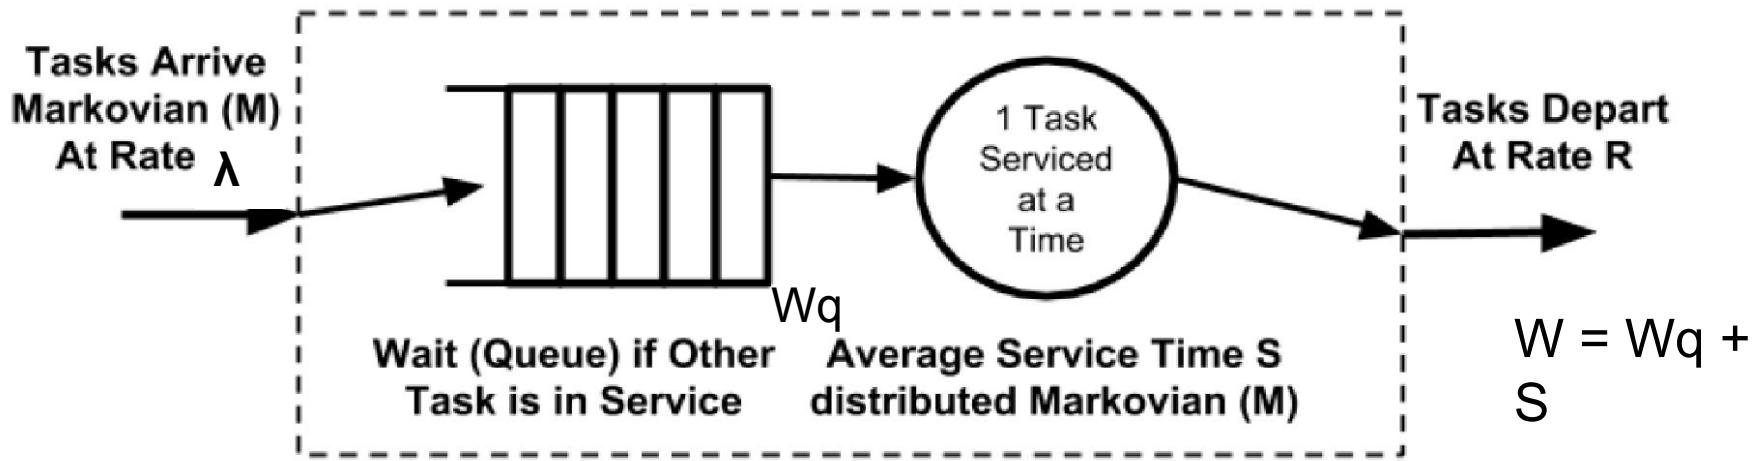
$$L = 0.3125 \text{ Gmisses/s} \times 100 \text{ ns} = 31.25 \text{ misses}$$

# M/M/1 Queue

$\text{M}$ : exponential interarrival times with mean  $1/\lambda$

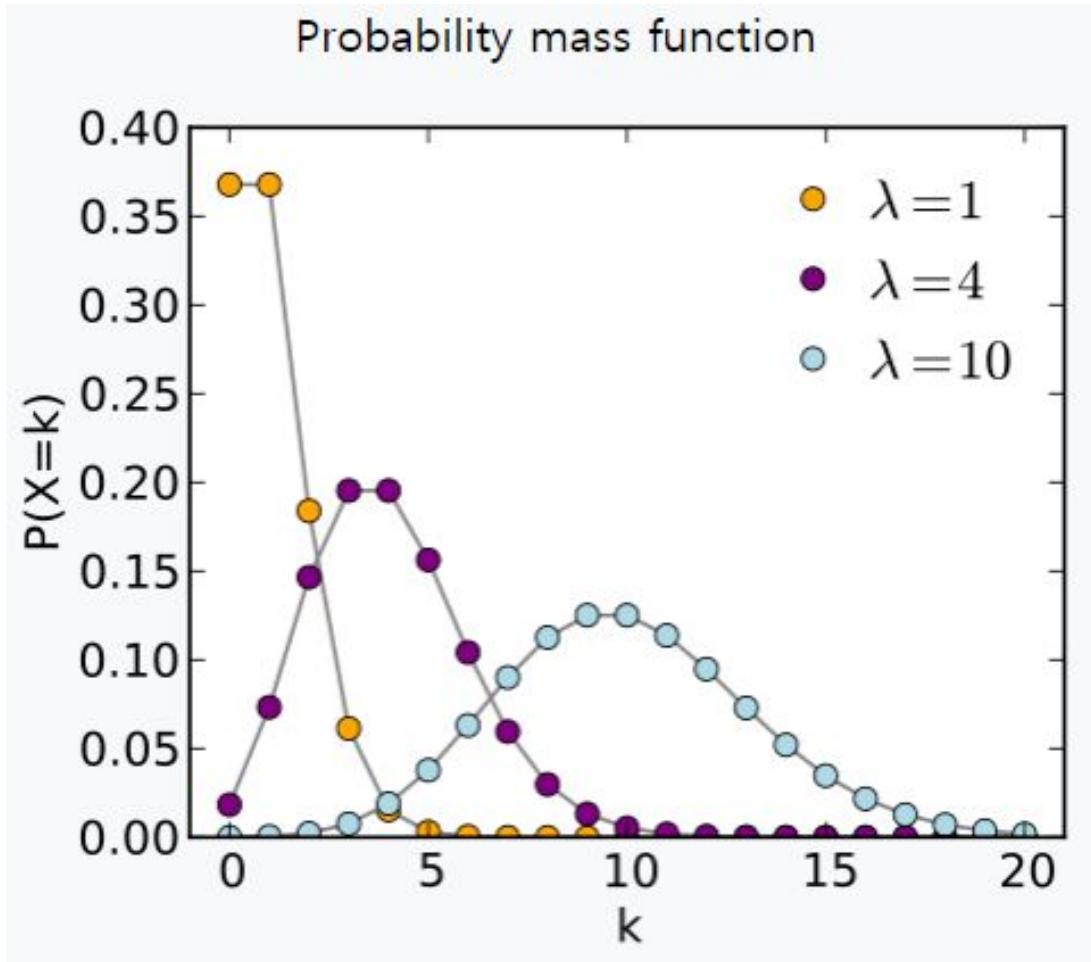
$\text{M}$ : exponential service times with mean  $S$

## 1. Single server



M/M/1 queue will reveal trade-off among  
(a) allowing unscheduled task arrivals, (b)  
minimizing latency ,  
and (c ) maximizing throughput

# Poisson Distribution



$k$ : the number of occurrences

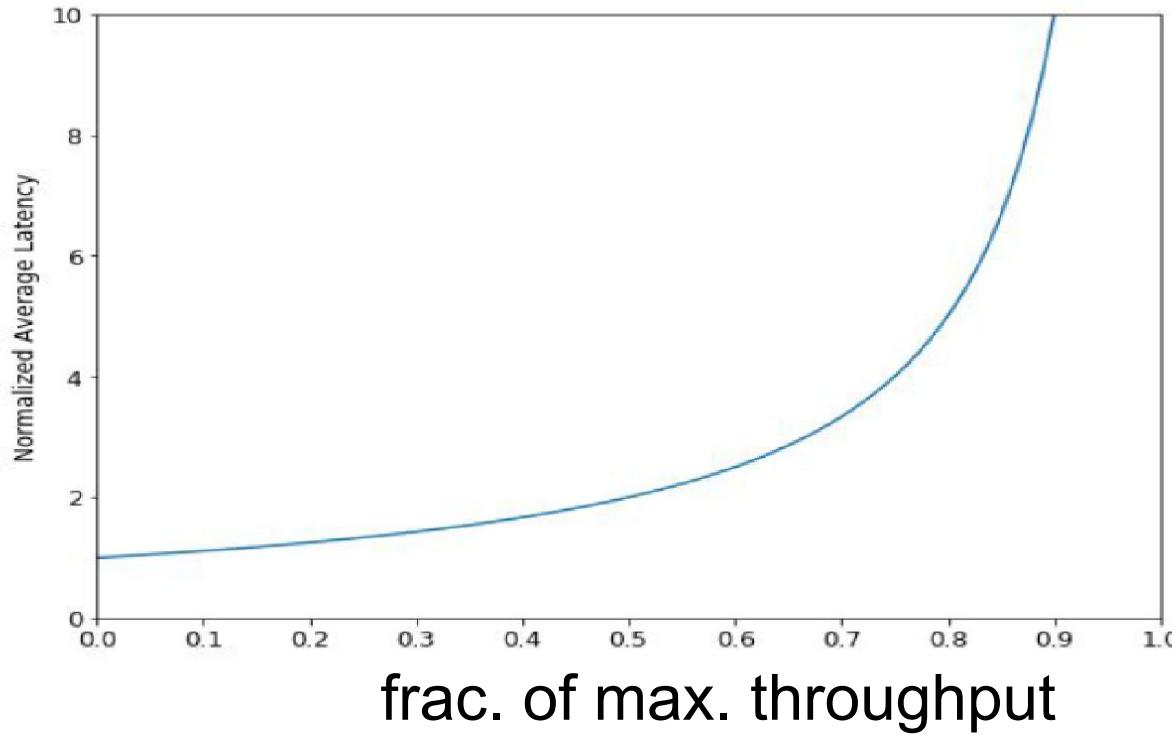
$\lambda$ : the expected rate of occurrences

[\[Wikipedia\]](#)

Performance Analysis Model].36

Jihong Kim (SNU)

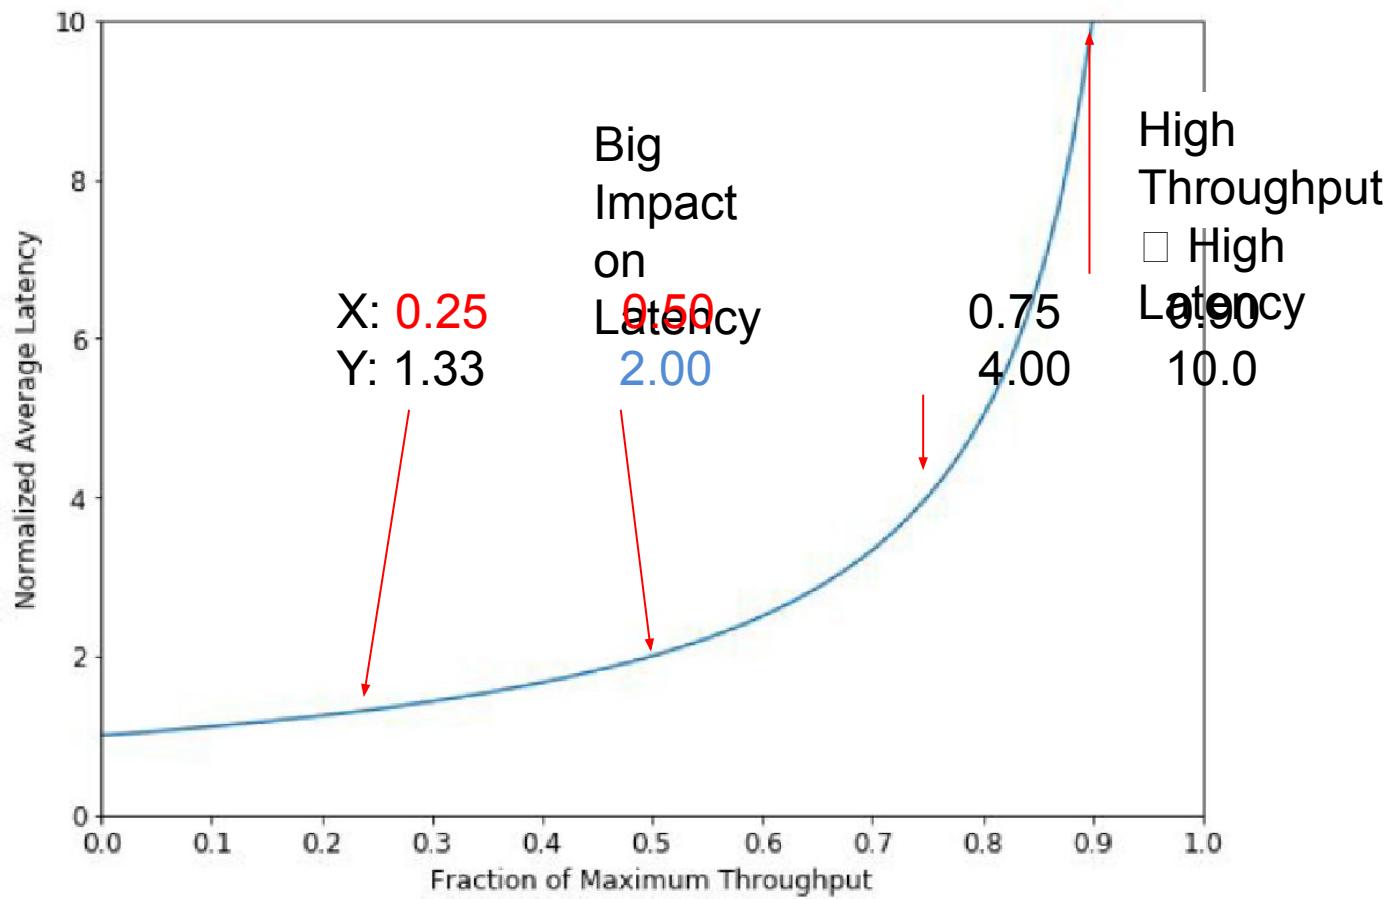
# M/M/1 Queue Latency vs. Throughput Tradeoff



From queuing theory,  $y = 1/(1-x)$

where  $y = W/S$  and  $x = \lambda S$

(utilization)



Low  
Throughput  
 Low  
Latency

# Example: M/M/1

Q: When  $\lambda = 10$  packets/s, if we want to keep  $W = 100$  ms,  
what  $S$  should be?

- Trial 1: Since  $1/\lambda = 100$  ms/packet, how about  $S = 100$  ms??
- In M/M/1,  $S = 50$  ms with  $Q = 50$  ms

(1) Solve for  $S$  in  $W/S = 1/(1 - \lambda S)$

$$1/10S = 1/(1 - 10S)$$

$$1 - 10S = 10S$$

$$S = 1/20 \text{ sec} \quad (\lambda \times S = 0.5 \square 50\% \text{ of max throughput})$$

(2) Solve for  $S$  in  $W = S/(1 - \lambda S)$

$$1/10 = S/(1 - 10S)$$

$$10S = 1 - 10S$$

$$S = 1/20 \text{ sec}$$

# Summary: Little's Law

- $L = \lambda \times W$  (items = throughput  $\times$  latency)
  - Assumptions:
    - stable, work-conserving; use long-run averages
- Design levers via  $L = \lambda W$ 
  - Cap L (WIP/concurrency)
  - Reduce W (service/overlap)
  - Bound  $\lambda$  (admission)

# Use Case: From SLO to Limits

- Given SLO  $W_{SLO}$  and expected  $\lambda$ ,
  - set  $L_{max} = \lambda \times W_{SLO}$
  - Enforce with max concurrent requests/tokens, queue depth
  - If  $L_{budget}$  is fixed,  $\lambda_{max} = L_{budget} / W$
  - Target utilization budget  $\rho$  between 0.6 and 0.8

# Example: LLM Inference Serving

- Inputs:  $\lambda = 100$  req/s, target  $W_{\text{SLO}} = 150$  ms
  - $L_{\text{max}} = 100 \times 0.15 = 15$   
 $\Rightarrow$  cap concurrent sequences  $\approx 15$
  - If current  $W = 120$  ms at  $\lambda = 100$   
 $\Rightarrow L \approx 12$  (headroom  $\sim 3$ )
  - Operate near  $\rho \approx 0.6\text{--}0.8$   
by tuning admission

# Example: NVMe Queue Depth for Checkpoints

- Checkpoint 20 GB every 120 s  $\Rightarrow \lambda \approx 171$  MB/s
  - Measured flush latency  $W \approx 400$  ms  $\Rightarrow L = \lambda W \approx 68$  MB in flight
  - With 1 MB I/Os  $\Rightarrow QD \approx 68$  (tune to p99, cap when saturation reached)
  - Re-evaluate if device S changes (e.g., thermal throttling, GC)

# M/M/1: Essentials & the Knee

- $\lambda$  arrivals,  $S = 1/\mu$  service time,  $\rho = \lambda S < 1$ 
  - $W = W_q + S$  (i.e., waiting + service)
  - $W = S/(1-\rho)$ ,  $W_q = \rho S/(1-\rho)$ ,  $L = \rho/(1-\rho)$
  - Rule-of-thumb:  
plan for  $\rho \approx 0.6\text{--}0.8$  to avoid latency blow-up
  - Check with  $L = \lambda W$  to keep units/budgets consistent

# References

- M. Hill, “Three Other Models of Computer System Performance,” 2019. Available at <https://arxiv.org/pdf/1901.02926.pdf>.
- S. Williams et al., “Roofline: An Insightful Visual Performance Model for Multicore Architectures,” CACM, April 2009.
- A. Ilic et al., “Cache-Aware Roofline Model: Upgrading the Loft,” IEEE CAL, 2014.
- A. Mohan, “Understanding Roofline Charts,” [http://telesens.co/2018/07/26/....](http://telesens.co/2018/07/26/)